

Numerical Optimization - Sheet 10

If you are a student in mathematics please solve the exercises with no tag and the ones with the tag **Mathematics**. If you are a data science student please solve the problems with no tag and those with the tag **Data Science**. Submissions with tags other than your subject count as bonus points. The tag **Programming** marks programming exercises.

Ex 1

(3 Points)

Let $f, g : \mathbb{R} \rightarrow \mathbb{R}$ be smooth functions. Consider the problem

$$\min_{\alpha \in \mathbb{R}} g(f(\alpha)) =: \min_{\alpha \in \mathbb{R}} F(\alpha).$$

- (i) The concatenation $g(f(\alpha))$ can be made implicit by the constraints $g(c)$ s.t. $f(\alpha) - c = 0$. Rewrite the given problem as constrained optimization problem without explicit concatenation (Don't forget the additional optimization parameter c).
- (ii) Compute the first order necessary optimality conditions (Theorem 2.24) of the constrained problem.
- (iii) Draw a connection between the adjoint λ in the optimal point and the derivative of F .

Solution 1:

(i)

$$\begin{aligned} \min_{\alpha, c \in \mathbb{R}} g(c) \\ \text{s.t. } c = f(\alpha). \end{aligned}$$

(ii) The lagrangian is of the form

$$\mathcal{L}(\alpha, c, \lambda) = g(c) + \lambda(f(\alpha) - c),$$

hence the KKT-Conditions state in particular, that

$$\nabla_{\lambda} \mathcal{L}(\hat{\alpha}, \hat{c}, \hat{\lambda}) = f(\hat{\alpha}) - \hat{c} = 0, \tag{1}$$

$$\nabla_c \mathcal{L}(\hat{\alpha}, \hat{c}, \hat{\lambda}) = g'(\hat{c}) - \hat{\lambda} = 0. \tag{2}$$

The derivative of F is given by the chain rule

$$F' = g'(f(\alpha))f'(\alpha).$$

Hence, in a stationary point we find $\hat{c} = f(\hat{\alpha})$ and therefore $\hat{\lambda} = g'(f(\hat{\alpha}))$, which is the first factor in the derivative of F .

Remark: In the language of deep neural networks we would say that the value $f(\alpha)$ is back-propagated. This propagation is clearly visible due to the adjoint λ in case of the constrained reformulation.

Ex 2 Programming

(2+2 Points)

Consider the example *Warm-up: numpy* on the pytorch tutorials page. The tutorial solves

$$\min_{a,b,c,d \in \mathbb{R}} \sum_{i=1}^N \|a + bx_i + cx_i^2 + dx_i^3 - \sin(x_i)\|_2^2, \quad (3)$$

for samples $x_i \in [-\pi, \pi]$ using a gradient descent method.

- (i) Propose a well known, classical method to solve the above problem in one step (you don't need to implement it).
- (ii) Change the tutorial's code such that it solves

$$\min_{a,b,c,d \in \mathbb{R}} \sum_{i=1}^N \|a + b \sigma(c + dx_i) - \text{sign}(x_i + 0.5)\|_2^2, \quad (4)$$

where $\sigma(x) = 1/(1 + \exp(-x))$ with $\sigma' = \sigma(1 - \sigma)$, and

$$\text{sign}(x) = \begin{cases} \frac{x}{|x|}, & x \neq 0 \\ 0, & \text{otherwise,} \end{cases}$$

on the same interval $x_i \in [-\pi, \pi]$.

Ex 3

(2+1+1 Points)

This exercise is a simple example for a minimization problem over a random variable where the stochastic gradient (SG) method with constant step size does not converge. To that end, let $f_r : [-1, 1] \rightarrow \mathbb{R}$ be defined as $x \mapsto (x + r)^2$ for $r \in [-1, 1]$. Consider the optimization problem

$$\min_{x \in [-1, 1]} \int_{-1}^1 f_r(x) \frac{1}{2} dr. \quad (5)$$

- (i) Compute the exact solution of problem (5) by hand.
- (ii) Show that for a constant step size $0 < \alpha \leq 0.5$ the iterates x^k of the SG-method do not leave the domain $[-1, 1]$.
- (iii) Show that for a constant step size $0 < \alpha \leq 0.5$ the algorithm does not converge to the solution.

Remark: The above problem can be interpreted as

$$\min_{x \in [-1, 1]} \mathbb{E} f(x)$$

if we assume that $r \sim \mathcal{U}[-1, 1]$ is uniformly distributed on $[-1, 1]$. In order to show, that the algorithm does not converge it suffices to consider the distribution of $x^+ = x^* - \alpha \nabla f_r(x^*)$ where x^* already is the solution.

Solution 3:

- (i) In this part we compute $f(x) := \mathbb{E}_r f_r(x)$. Note, that in practice this function is not known, and

the stochastic gradient method does not depend on the knowledge of this function.

$$\begin{aligned} f(x) &= \int_{-1}^1 f_r(x) \frac{1}{2} dr = \int_{-1}^1 (x+r)^2 \frac{1}{2} dr \\ &= \int_{-1}^1 (x^2 + 2xr + r^2) \frac{1}{2} dr = x^2 + \left[\frac{2}{4} x r^2 \right]_{-1}^1 + \left[\frac{1}{6} r^3 \right]_{-1}^1 \\ &= x^2 + \frac{1}{3}. \end{aligned}$$

Now as we know f we can compute the solution of problem (5), which is $x^* = 0$.

- (ii) We now want to apply the SG-method. Before we can do so we make (a rather technical) observation, namely: *If we apply the SG-method the iterates x_k stay in $[-1, 1]$.*

The SG-method for parameter x and a random function $f_r(x)$ is given by

```
choose a starting value x = x0
choose a learning rate alpha

for k=1,...,maxit:
    choose a random number r
    evaluate g_r = gradient(f_r(x))
    update x = x - alpha * g_r
```

The derivative of f_r at x is given by

$$g(x) = 2(x + r).$$

Hence, if $x_1 \in [-1, 1]$ we obtain via induction that

$$|x_{k+1}| = |x_k - 2\alpha(x_k + r_k)| = |x_k(1 - 2\alpha) + 2\alpha r_k| \leq |x_k|(1 - 2\alpha) + 2\alpha|r_k| \leq 1.$$

This means that we can safely apply the SG-method as the iterate x_k will not leave the domain $[-1, 1]$.

- (iii) As given in the hint we only consider the distribution of $x_{k+1} = x_k - \alpha \nabla f_r(x_k)$, where $x_k = x^* = 0$ is already the solution. We show now that the variance $\mathbb{E}\|x^* - x^+\|^2$ is positive and that it does not depend on k , so that there is no convergence.

As $x^* = 0$ and we assume that $x_k = 0$ we get that

$$\mathbb{E}\|x^* - x_{k+1}\|^2 = \mathbb{E}\|x^* - (x_k - 2\alpha(x_k + r))\|^2 = \mathbb{E}\|2\alpha r\|^2,$$

where the term on the left does not depend on k anymore. Finally we can compute

$$\begin{aligned} \mathbb{E}\|2\alpha r\|^2 &= \int_{-1}^1 (2\alpha r)^2 \frac{1}{2} dr = \int_{-1}^1 4r^2 \alpha^2 \frac{1}{2} dr \\ &= 2\alpha^2 \left[\frac{1}{3} r^3 \right]_{-1}^1 = \frac{4\alpha^2}{3} > 0, \end{aligned}$$

which again does not depend on k . Hence, there is no convergence, and we would observe an oscillating behavior even when the algorithm gets close to the solution. The oscillation would be smaller for smaller α but it cannot vanish.

Ex 4

(2* Bonus- Points)

Assume you are given an input sample $x^{(i)} \in \mathbb{R}^n$ which could be an image or sound etc. which you feed into a trained neural network $N : \mathbb{R}^n \rightarrow \mathbb{R}^d$. So we interpret the network as function which maps an input to a probability distribution on $\{1, \dots, d\}$ and forget its parameterization for now. The input sample $x^{(i)} \in \mathbb{R}^n$ has a true class, namely $i \in \{1, \dots, d\}$. Assume for now that the network maps the sample to its correct class, i.e. assigns the highest probability to class i . The following optimization problem aims to find a distortion of $x^{(i)}$ by a small vector $\Delta x \in \mathbb{R}^n$.

Let $0 < \varepsilon < 1$, and $k, i \in \{1, \dots, d\}$, $k \neq i$. Let moreover $x^{(i)} \in \mathbb{R}^n$ be a fixed input sample of class i and $N_j(x)$ be the j -th component of $N(x)$ for $j \in \{1, \dots, d\}$. Consider the problem

$$\begin{aligned} & \min_{\Delta x \in \mathbb{R}^n} \|\Delta x\|_2^2 \\ & \text{s.t. for all } j \in \{1, \dots, d\}, j \neq k \\ & N_k(x^{(i)} + \Delta x) \geq (1 + \varepsilon) N_j(x^{(i)} + \Delta x). \end{aligned}$$

What is the purpose of the given problem? What would be the consequence of the accessibility of very good solutions to it?

Solution 4:

- (i) The Problem tries to find the minimal distortion of $x^{(i)}$ such that the network classifies it as class $k \neq i$. A solution $x^{(i)} + \Delta x$ is called (targeted) adversarial example.
- (ii) It seems that many network architectures are prone to adversarials which are that small, that they are unrecognizable for humans. Hence a foto or voice can look or sound unchanged while being recognized as totally different for the network.